

Point cloud method for rendering BSSRDFs

Alessandro Dal Corso
Technical University of Denmark

Jeppe Revall Frisvad
Technical University of Denmark

March 2018

This is a short note on rendering a triangular mesh onto which we apply a BSSRDF model. The method works on arbitrary static triangular meshes and is unbiased.

We first define some quantities relative to our mesh. We consider a triangle mesh a set $M = \{T_i, i \in [0, N_\Delta - 1]\}$ composed of N_Δ triangles. Each triangle T_i is composed of three vertices:

$$T_i = \{\mathbf{v}_0^i, \mathbf{v}_1^i, \mathbf{v}_2^i\}$$

From these quantities it is straightforward to define two derived quantities, the normal \vec{n}_i and the area A_i of the triangle.

$$\vec{n}_i = \frac{(\mathbf{v}_1^i - \mathbf{v}_0^i) \times (\mathbf{v}_2^i - \mathbf{v}_0^i)}{\|(\mathbf{v}_1^i - \mathbf{v}_0^i) \times (\mathbf{v}_2^i - \mathbf{v}_0^i)\|}$$
$$A_i = \frac{1}{2} \|(\mathbf{v}_1^i - \mathbf{v}_0^i) \times (\mathbf{v}_2^i - \mathbf{v}_0^i)\|$$

Once we have defined our triangles, we can start describing our solution.

Theoretically, we solve the standard reflectance equation for BSSRDFs:

$$L_r(\mathbf{x}_o, \vec{\omega}_o) = \int_A \int_\Omega S(\mathbf{x}_i, \vec{\omega}_i, \mathbf{x}_o, \vec{\omega}_o) L_i(\mathbf{x}_i, \vec{\omega}_i) (\vec{n}_i \cdot \vec{\omega}_i) dA_i d\vec{\omega}_i$$

We now solve the integral with Monte Carlo integration with importance sampling. We take K samples from the area integral and one sample from the directional integral:

$$\begin{aligned} \hat{L}_r(\mathbf{x}_o, \vec{\omega}_o) &\approx \sum_{k=1}^K \frac{S(\mathbf{x}_{i,k}, \vec{\omega}_{i,k}, \mathbf{x}_o, \vec{\omega}_o) L_i(\mathbf{x}_{i,k}, \vec{\omega}_{i,k}) (\vec{n}_{i,k} \cdot \vec{\omega}_{i,k})}{\text{pdf}(\vec{\omega}_{i,k}) \text{pdf}(\mathbf{x}_{i,k})} = \\ &= \sum_{k=1}^K S(\mathbf{x}_{i,k}, \vec{\omega}_{i,k}, \mathbf{x}_o, \vec{\omega}_o) \Phi_{i,k}(\mathbf{x}_{i,k}, \vec{\omega}_{i,k}) \end{aligned} \tag{1}$$

The algorithm is composed of two phases. In the first phase, we generate the K samples on the surface of the model, storing for each position $\mathbf{x}_{i,k}$, direction $\vec{\omega}_{i,k}$ and incoming flux $\Phi_{i,k}$. In the second phase, we render the final image evaluating the contribution at each pixel.




Figure 1: Diagram illustrating our method. Two points $\mathbf{x}_{i,k}$ are selected on the surface, and the light stored in a buffer. The contribution is then gathered from the various points out towards the eye at point \mathbf{x}_o .

In phase one, we produce K samples. The k -th sample is composed of a position $\mathbf{x}_{i,k}$, a direction $\vec{\omega}_{i,k}$ and a flux $\Phi_{i,k}$. First, we choose a triangle T_j . We choose the triangle uniformly, so given a random number $\xi \in [0, 1)$, $j = \lfloor \xi N_\Delta \rfloor$. To get \mathbf{x}_k , we sample a point uniformly on the triangle T_j :

$$\mathbf{x}_{i,k} = (1 - \sqrt{\xi_0})\mathbf{v}_0^j + (1 - \xi_1)\sqrt{\xi_0}\mathbf{v}_1^j + \xi_1\sqrt{\xi_0}\mathbf{v}_2^j$$

Where $\xi_0, \xi_1 \in [0, 1)$.

Now, the sampling direction $\vec{\omega}_{i,k}$ and the flux $\Phi_{i,k}$ need to be evaluated. We leave the choice of the sampling distribution for $\vec{\omega}_{i,k}$ to implementation, since it is not important for the method. From sampling the light source we obtain $\vec{\omega}_{i,k}$ and an irradiance $E_{i,k}$:

$$E_{i,k}(\mathbf{x}_{i,k}, \vec{\omega}_{i,k}) = \frac{L_i(\mathbf{x}_{i,k}, \vec{\omega}_{i,k})(\vec{n}_{i,k} \cdot \vec{\omega}_{i,k})}{\text{pdf}(\vec{\omega}_{i,k})}$$

We need the pdf of choosing the point $\text{pdf}(\mathbf{x}_{i,k})$. Since we sampled the triangle and then the point, we combine the two independent probabilities:

$$\text{pdf}(\mathbf{x}_{i,k}) = \frac{1}{N_\Delta A_j}$$

Putting it all together:

$$\Phi_{i,k}(\mathbf{x}_{i,k}, \vec{\omega}_{i,k}) = \frac{E_{i,k}(\mathbf{x}_{i,k}, \vec{\omega}_{i,k})}{\text{pdf}(\mathbf{x}_{i,k})} = N_\Delta A_j E_{i,k}(\mathbf{x}_{i,k}, \vec{\omega}_{i,k})$$

Once we have stored this quantity in the sample, we are ready to proceed to the second and final phase. For each pixel, we ray trace a camera ray, obtaining a point \mathbf{x}_o and a direction towards the camera $\vec{\omega}_o$. We consider the scattering medium with coefficients σ_a, σ_s, g and relative index of refraction η as a pure Fresnel interface with added scattering.

In our path tracing implementation, we perform russian roulette on the Fresnel coefficient R to choose reflected or refracted direction. In the case of reflection, we continue tracing the reflected ray. In the case of refraction, we calculate the final radiance as:

$$\hat{L}_o(\mathbf{x}_o, \vec{\omega}_o) = L_e(\mathbf{x}_o, \vec{\omega}_o) + L_t(\mathbf{x}_o, \vec{\omega}_o) + \hat{L}_r(\mathbf{x}_o, \vec{\omega}_o)$$

Where L_e is the emitted radiance from the medium, L_t is the radiance due to direct transmission and \hat{L}_r comes from 1.

To efficiently render the model and save expensive BSSRDF evaluations, we do a last optimization. We evaluate the BSSRDF stochastically with probability $\exp(-\sigma_{tr}\|\mathbf{x}_o - \mathbf{x}_k\|)$, where $\sigma_{tr} = \sqrt{3\sigma_a(\sigma_a + (1-g)\sigma_s)}$ is the effective transport coefficient. So, only samples that are very close to the exit point \mathbf{x}_o will be actually evaluated. In formulas:

$$\hat{L}_r(\mathbf{x}_o, \vec{\omega}_o) \approx \sum_{k=1}^K S(\mathbf{x}_k, \vec{\omega}_k, \mathbf{x}_o, \vec{\omega}_o) \Phi_{i,k}(\mathbf{x}_k, \vec{\omega}_k) V(\xi, e^{-\sigma_{tr}\|\mathbf{x}_o - \mathbf{x}_k\|}) e^{\sigma_{tr}\|\mathbf{x}_o - \mathbf{x}_k\|}$$

Where:

$$V(\xi, d) = \begin{cases} 1 & \text{if } \xi < d \\ 0 & \text{otherwise} \end{cases}$$

In case of a spectral RGB rendering, we pick σ_{tr} as the mean of the three rgb components.